

GRS

Management vs. Operations

- **Management:**
 - Conceptual, high-level dimensions to keep network running smoothly;
 - **Smoothly:** Quality of Experience (QoE), reliability, security, etc.
 - Manage monitoring, configurations, performance, faults, security, accounting, etc.
- **Operations:**
 - People, processes, and tools to make management happen; Typically in the network operations center.

Dimensions of networking

Packets

- **Data plane:**
 - Packet switch and forwarding (microseconds);
 - Looks up rules, e.g., packet header, in/out interfaces, time-slot;
- **Control plane:**
 - **Switching control/Network functions** (milliseconds): routing, Quality of Service (QoS), firewall, etc.
 - Reacts to changes (interfaces, links, etc.);
 - Control is centralized or distributed (neighboring routers);
 - Update switching rules.
- **Management plane:**
 - Orquestração e operação de redes;
 - Protocolos: Telnet, SSH, SNMP.
- **Network services** (DHCP, DNS, Radius, LDAP, etc.) (milliseconds):
 - Provide mapping: MAC-IP, IP-Domain name, accounting, etc.

Management

- Configuration - Provisioning, software updates, structural changes, minor configuration changes;
- Performance assessment;
- Accountability;
- Fault management;
- Security.

Changing network topologies

- **Spanning Tree** - focus on Ethernet reliability while avoiding loops;
- **Fat tree** - focus on bandwidth between nodes;
- Clos network, crossbar switch, multistage:
 - Strict-sense nonblocking: unused input => unused output;
 - Rearrangeably nonblocking: need to rearrange the connections between middle switches.
- Traffic demands:
 - North-south (computer to Internet), typical enterprise - spanning tree;
 - East-west (computer to computer), data center - fat tree.

Services and systems

- Traditional network-based services:
 - Mail, web, storage (NAS/SAN), database, etc.
- **Systems** - more complex services, e.g.:
 - Three tier applications: database backends, user-facing applications, application;
 - Micro-services, REST APIs, publish-subscribe;
 - Big data.

Devops in the cloud

- **Traditional** release deployment:
 - Gather specifications;

- UML architecture diagrams;
- Implement, test, deploy;
- **Fence between *dev team* and *ops team*;**
- Dev not consider operational requirements;
- Broken deployments => long feedback to dev.
- **Devops** approach:
 - Quick deployment cycle (agile, test-driven, sprints);
 - Write code thinking about other phases (deployment, testing, etc.);
 - Build rather than buy;
 - Automate test and deployment (repeatable and predictable);
 - Embrace failure: fail fast, find errors/vulnerabilities, recover quickly.

Devops for networking

- Cloud environment easier because it's virtual;
- Network is *hard*:
 - Physical limitations (e.g., cables, device access);
 - Hard to control (heterogeneous device API);
 - Complex (multiple protocols, topologies);
 - Hardware bundle (switching, routing, etc. in the same closed box).
- If possible:
 - Quick deployment;
 - Predictable deployment;
 - Fail fast;
 - Build rather than buy.

How to

- Virtual networks;
- Software-defined networks (can use SDN);
- Network function virtualization (enables running network functions in generic hardware);
- Programmatic device interfaces;
- Anyone can write code.

FCAPS

Framework for network management.

Component	Explanation
Fault	Recognize, isolate, correct and log faults in the network
Configuration	Focus tracking and managing deployments in a centralized manner. Streamline device configuration and provisioning to ensure quicker configuration and flexibility
Accounting	gather usage statistics for users
Performance	ensuring performance remains at acceptable levels
Security	controlling access to assets in the network

Monitoring entra em todas exceto configuration.

Network Operation Center (NOC)

Centralized location where IT teams can continuously monitor the performance and health of a network.

- Serves as the first line of defense against network disruptions and failures;
- Through the NOC, organizations gain full visibility into their network;
- Detect anomalies and either take steps to prevent problems or quickly resolve issues as they emerge;
- Can be built internally and located on-premise (often within the data center).

NOC Tools

- Incident tracking through a **Ticketing System**;
- Centralization of knowledge;
- Daily and monthly reports for measuring incident severity and improvement progress;
- Monitoring - NOC has 2 monitoring process types that are relevant to it: infrastructure monitoring + user experience monitoring;
- Automation of the IT process.

Monolithic vs. microservice

Monolithic

All functionalities of a project exist in a single codebase.

Good	Bad
Simpler to develop	Becomes too large with time => difficult to manage
Easier to deploy	Need to redeploy the whole application with every change
Network latency and security not as relevant	High barrier of entry to new devs
Developers focus on single application	Even if single part is suffering load, need to replicate instances for the entire application
–	Difficult to adopt new techs
–	Not very reliable as single bug can bring everything down

Microservices

Architectural development style where the application is made of smaller services. These communicate with each other directly using lightweight protocols like HTTP.

Good	Bad
Single responsibility	More complex (distributed system)
Built around business capabilities - freedom to choose the best technologies for each business part	Independent deployment of services is hard
Design for failure - designed with failure cases in mind. A single service failing, doesn't bring everything down	Debugging is difficult

Virtualization vs. containers

Virtualization	Containers
Needs hypervisor	Shares the OS kernel
Brings a whole system (OS and virtual hardware)	Shared parts are read-only
Moving VMs between clouds can be challenging	Doesn't allow running other OSes

Quality Network

Challenges

- Large number of devices + switching limitations:
 - Capacity bottlenecks.
- Specific requirements for users and applications;
- Costs;
- Traffic growth;
- Network outages;
- Users with different levels of access and accounts;
- Attacks and security.

Property	Virtualization	Containerization
Isolation	Fully isolates the host operating system and virtual machines	Isolates the host and other containers to a certain degree; doesn't provide a strong security boundary between hosts and containers
Operating System	Includes a separate, completely independent operating system with the kernel and requires more CPU, memory, and storage resources	Involves a user-mode operating system that can be tailored to contain only those services your app needs, so its light on resource requirement
Guest Compatibility	Compatible with almost all operating systems inside a virtual machine	Compatible only with a similar operating system version as the host
Deployment	Can be deployed individually with a hypervisor for each VM	Deploys individual containers through Docker and multiple containers with Kubernetes orchestration
Persistent Storage	Uses a virtual hard disk (VHD) for single VM local storage or server message block (SMB) for shared storage on multiple servers	Uses local disks for local storage for a single node and SMB for shared storage on multiple servers or nodes
Load Balancing	Runs VMs on other servers in a failover cluster for load balancing	Manages load balancing by automatically starting and stopping containers on cluster nodes through an orchestrator
Networking	Conducted via virtual network adapters (VNA)	Uses an isolated view of a VNA for lightweight virtualization

Figure 1: virtualization vs. containers

Network design

- How traffic flows: north-south, east-west;
- Different parts of the network may use different technologies and topologies: LAN, WAN, MAN, STP, fat tree, etc., speed;
- Segregation: workstations, servers/datacenter, public-facing (DMZ), admin, etc.;
- Technologies: Ethernet, VLAN, EtherChannel, MPLS, IP, OSPF, BGP;
- Interconnection with other networks: ISP, BGP peers, other networks of other departments.
- **DMZ** - Expose part of the network to the public without isolating that part from the rest of the network;
- **Point of Presence (PoP)** - local access point for an ISP:
 - Consists of high-speed equipment;
 - Enables users to connect to the internet through the ISP;
 - One or more unique IPs + pool of assignable IPs.
- **Internet Exchange Points (IXPs)** - physical location through which Internet infrastructure companies such as ISPs and CDNs connect with each other:
 - Exist at the edge of different networks;
 - Allow network providers to share traffic outside their own networks;
 - Reducing latency, improving round-trip time, and potentially reducing costs.

The “dev | fence | ops” trap configuration

- **Silos:**
 - Network planning and design;
 - Network deployment.
- **Devops** for networking allows a more iterative process of design, deploy, and getting feedback to update the network design;
- **Network Function Virtualization (NFV) => helps;**
- **Cost of hardware and hardware compatibility with future network expansions => hinders.**

Application quality requirements

- Capacity (bit/s) - bandwidth intensive applications:
 - Bursts: o quão frequentes e de quão grandes;
 - Capacity vs. throughput vs. goodput.
- Delay - real time, interactive applications:
 - End-to-end delay (control);
 - Round-trip delay (teleconference);
 - Delay variation/jitter - visualization.
- Reliability - mission critical applications:
 - Error rates nos bit/packets;
 - Mean time between failures (MTBF);
 - Mean time to recover (MTTR);
 - Availability = $MTBF / (MTBF + MTTR)$;
 - Uptime(%) = 1 - Availability;

QoS and traffic engineering

- **Best effort networks:**
 - Não tem preocupações de qualidade (e.g., queues);
 - Serve pacotes por ordem de chegada;
 - Light used => quality OK;
 - Heavily used => quality degradation.
- **QoS:**
 - Queue management - choose packet, different queues;
 - Load-balancing algorithms - Round-robin, token bucket, RED, etc.
- **Circuit Switching Networks:**
 - Bandwidth is divided into chunks;
 - The Bit delay is constant (during the connection);
 - Once path/circuit is established, data rate is constant;

- **Admission control** - validation step in communication systems to see if there are sufficient resources for the connection.

Security

- Enforcing security:
 - Segregation;
 - Access control;
 - Firewall;
 - IDS/IPS.
- Security management:
 - Vulnerability scanning;
 - Intelligence gathering;
 - Incident response;
 - Forensics.

Monitoring

- Get a sense of how the network is performing;
- Make sure we're still offering a quality network;
- Essential for **FCAPS**;
- **Short-term measurements**: identify faults, congestions, and attacks;
- **Longer-term measurements**: traffic engineering (reroute or negotiate new agreements with peers), upgrade link and device capacity;
- **Accounting** - so you know how each client of the network is using the network.

Types of measurements

- Application and user-related measurements. E.g., web page loading time;
- Device measurements. E.g., CPU, memory, disk, temperature;
- Network measurements. E.g., traffic data (packet traces, flow data), latency, throughput.

Passive vs. Active network measurements

- **Passive**:
 - Get a sense of the existing traffic in the network;
 - Have devices report how much traffic is going through (SNMP, netflow);
 - **Port mirroring** - Tap a link or copy traffic to monitoring port;
 - For measuring production traffic and its characteristics;
 - **Packet traces** - Motion-picture-like of everything that goes through the network. Raw data is powerful but hard to use. Difficult to manage, use, and process;
 - **Traffic counters** - Routers keep track of how much traffic goes through each link. Simple but limited in scope;
 - **Flow measurements** - keeps record of traffic for each flow. Simpler to use. More info. Each operation generates many TCP flows.
- **Active**:
 - Inject (new) measurement packets in the network;
 - Get a sense of how the network reacts to these packets;
 - Including responses (e.g., ICMP req/rep for RTT);
 - For measuring the properties of the network (delay, jitter, topology, etc.);
 - **Tools**: ping + traceroute + lperf + owamp + twamp.
- **Não precisamos de ferramentas de medidas passivas para fazer medidas ativas.**

What to do with measurement data

- Store for later query and processing:
 - Send to ELK (big data storage);
 - Plot charts, and queries on past data;
 - Build historical dataset for learning AI models for different management tasks.
- Use immediately once data is generated:
 - Anomaly detection;

- Apply state rules, use pretrained AI;
- Online learning (update AI).

Traffic engineering

Apply when congestion arises: **Admission Control** + policing and shaping + queuing and scheduling policies.

- **On the device:**
 - Token bucket, fair queuing, etc.
- **On the network:**
 - ATM (dead), IP (intserv), label switching (diffserv, MPLS/IP).
- For protection:
 - Pre-routed alternative paths;
 - Don't wait until link failure to start routing algorithm;
 - Start sending through alternative route when link failure detected.
- For load balancing:
 - Multipath TCP, DNS, reverse proxy.

IP Addressing

Classes

Class	Format
A	/8
B	/16
C	/24
Classless (CIDR)	Specify n bits in network part (/n)

Special cases

Meaning	Address
Network address	host part all 0s
Broadcast address	Host part all 1s
This network, this host	0.0.0.0/8
Loopback	127.0.0.0/8
Link local	169.254.0.0/16, FE80:0:0:0:<Interface ID>
Private, NAT	10.0.0.0/8, 172.16.0.0/12, 192.168.0.0/16

Splitting networks

Example: split 200.17.30.0/24 ($2^{32-24} = 2^8 = 256$ addresses).

- How many 32 address networks fit in this /24 network?
- $32 = 2^5 \rightarrow 5$ bits \Rightarrow /27 network;
- $8 - 5 = 3$ bits for subnetting $\Rightarrow 2^3 = 8$ subnets;
- 200.17.30.xxyyyyyy, where x=subnet, y=host.

Aggregating networks

Example: is it possible to aggregate any pair of /27 into one /26?

- 2nd and 3rd /27 subnets;
- 2nd: 200.17.30.32 to 200.17.30.63 - 200.17.30.00|1yyyyy
- 3rd: 200.17.30.64 to 200.17.30.91 - 200.17.30.01|0yyyyy
- Even though it is continuous, it doesn't work.

How to?

- Networks must be contiguous: no missing addresses between first and last addr;
- All addrs in same network must have same **network ID bits**.

DNS

2 parts

- Establish the map between domain names and IPs;
- Answer queries from users:
 - Can ask the authoritative name server;
 - User can group together so that responses can be cached. The burden can be eased on authoritative servers (cache, proxy, resolver DNS)

Domain name hierarchy

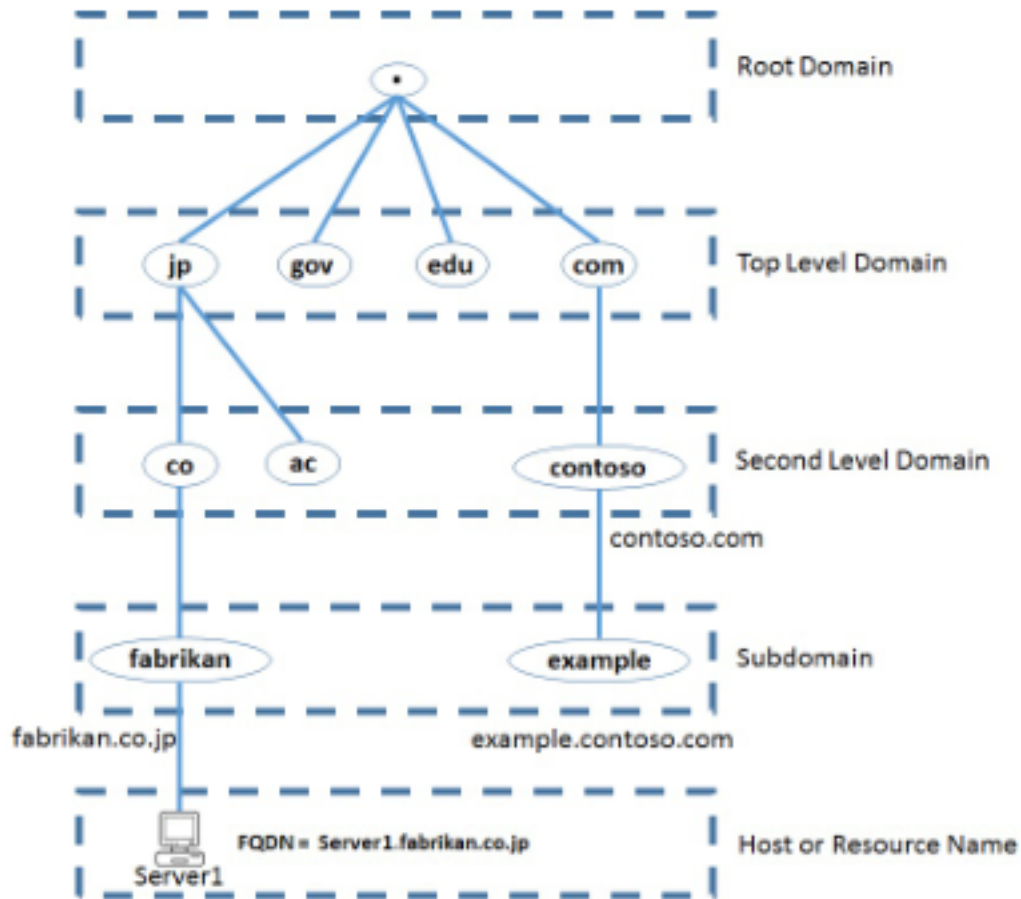


Figure 2: domain name hierarchy

Domain name specifications

- Entities:
 - Internet Corporation for Assigned Names and Numbers (ICANN)
 - Internet Assigned Numbers Authority (IANA)
- Format:
 - ASCII (original)
 - Internationalized Domain Names - multibyte unicode, encoded in ASCII (punycode)

Root Authoritative Servers

- **DNS zone** - portion of the DNS namespace that is managed by a specific organization or administrator;
- **DNS root zone** - contains all domain names;
- Root servers respond to TLD requests.

Non-root Authoritative Servers

- **Top-Level Domains:**
 - gTLD - Generic Top-Level Domains;
 - ccTLD - Country Code Top-Level Domains.
- **Second-Level Domains:**
 - Each TLD has a TLD manager organizations - responsible for assigning SLDs under that TLD;
 - The same happens down the hierarchy until the final host or resource name.
- **Masters:**
 - Stores the definitive versions of all records;
 - Identified in start-of-authority (SOA) resource record.
- **Slaves:**
 - Automatic updating mechanism to maintain an identical copy of the primary server's database for a zone;
 - DNS zone transfer (AXFR).

Architecture and Protocol

- **Local:**
 - /etc/hosts, resolv.conf, DHCP;
 - Non-recursive, recursive, iterative.
- **Servers:**
 - Local/remote DNS server/cache/resolver;
 - Root name server, TLD, SLD;
 - Caching;
 - Split server.

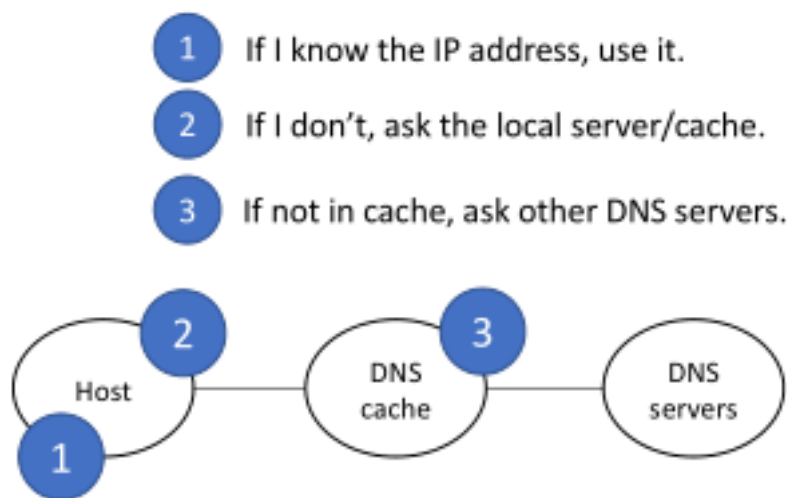


Figure 3: dns_cat_ask_order

DNS message format

- QR (1 bit). 0 => query. 1 => response.
- AA (1 bit). 1 is authoritative
- Query type - what thing we want (A, CNAME, MX, etc.). For **queries**.
- type - the type it is answering. For **responses**. Matches the **query type**.
- Query name - Numero de chars (0-63) de cada level e depois o level. Acaba em 0.
- Query class - 1 for IPs.

Timeouts

- **Expire** - number of seconds after which **secondary name servers** should stop answering requests for this zone **if the master does not respond**:
 - Should be higher than *Refresh* + *Retry*.
- **Refresh** - period of refreshing;

NS OVER TOR, ...

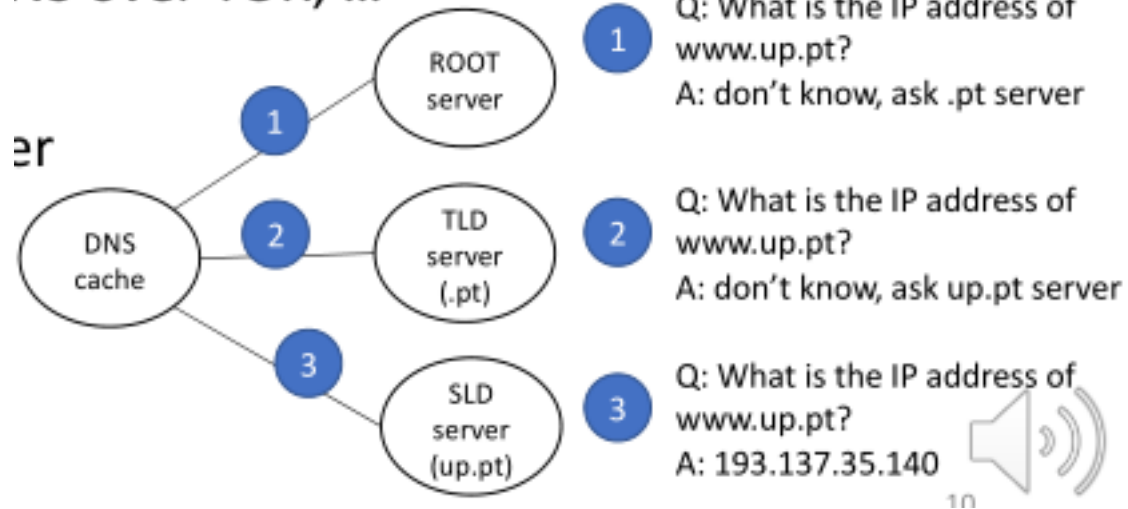


Figure 4: server_hierarchy




Figure 5: domain_name_to_query

- **Retry** - quanto tempo pode estar sem conseguir dar refresh até dizer “primary server failure”.

Resolvers and caches

- Não tem propriamente load-balancing:
 - Cliente tem info de primários + secundários;
 - Pergunta ao primário e depois aos secundários por ordem quando as coisas falham.

Resolvers and caches

• Recursive

vs.

• Iterative

• Referral

- NS: “can’t find this name, try this NS”
- (2) and (4) in red in the figure

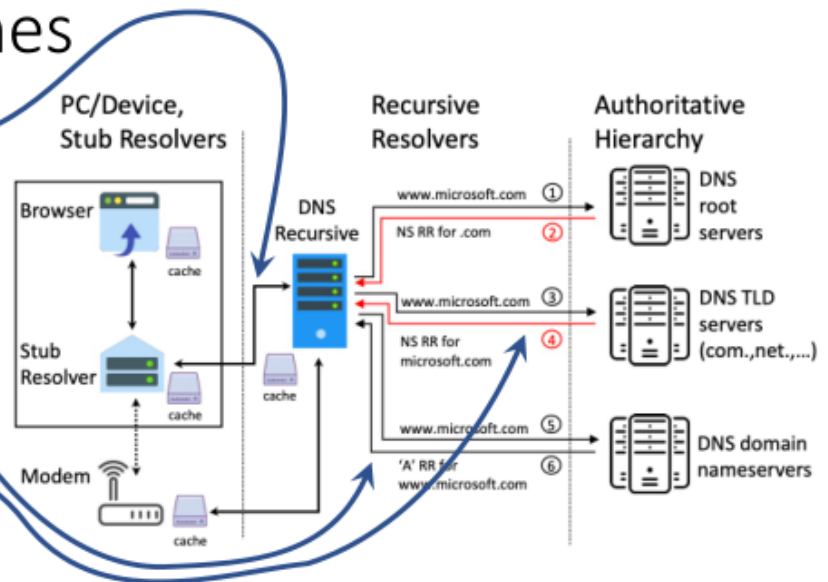


Figure 1: The resolution process, in theory, for the domain `www.microsoft.com`. The red steps represent NS referral responses.

Figure 6: Resolvers and Caches

Transport Protocols for DNS

- Typically UDP port 53;
- Can be TCP port 53:
 - Large replies (UDP max 512 bytes);
 - Zone transfers (AXFR).
- Encrypted:
 - DoT; DoH; DNS over TOR.
 - Privacy;
 - Easy to recover domain name of HTTP request from DNS query, even if HTTP and Server Name Indicator encrypted.
 - DNSSEC only gives integrity of query answers.

Routing

Inside a router

- Each router has its own routing table;
- If routing tables aren't coherent => loops, blackholes.
- Routers can sense if their links are up or down (or if new networks are added);
- Automatically update routing tables, based on some metric to optimize:
 - Smallest number of hops;
 - Smallest delay;
 - Largest throughput;
 - Balance load on links;
 - Administrative distance.




Figure 7: routers_routing

Centralized routing

- Nodes report link status to central location;
- Central location runs optimization algorithm;
- Return routing table to each router;
- **Uses => SDN-like solution** - data centers.

Distributed routing

- Update routing tables dynamically;
- According to information from other nodes:
 - Each router knows the networks it is directly connected to;
 - Routers send control packets to other routers.
- **Uses => Geographically dispersed network.**

Routing algorithms

- **Link state:**
 - Full topology of the network then Dijkstra;
 - Distributed: flood link state announcements - OSPF.
- **Distance vector:**
 - Distance to destination node, Bellman-Ford;
 - Distributed: update distances, resend to neighbors.
- **Path vector:**
 - List of nodes to destination node, local preferences;
 - Distributed: update path vectors, resend to neighbors.

Interior routing

- Same administrative domain:
 - All networks and devices administered by same entity;
 - Typically the internal network of an organization.
- Protocols:
 - RIP (outdated) - distance vector;
 - EIGRP (Cisco) - distance vector, optimized, hybrid;
 - **OSPF - link state.**

OSPF

- Messages:
 - Hello packets between routers to establish adjacency;
 - Link state advertisement, link-state database.
- Designated routers, network adjacency vs. point-to-point link;
- Cost metric manually defined, or multiple of reference bandwidth;
- Backbone area, separate link-state database, Area Border Router.

Exterior routing

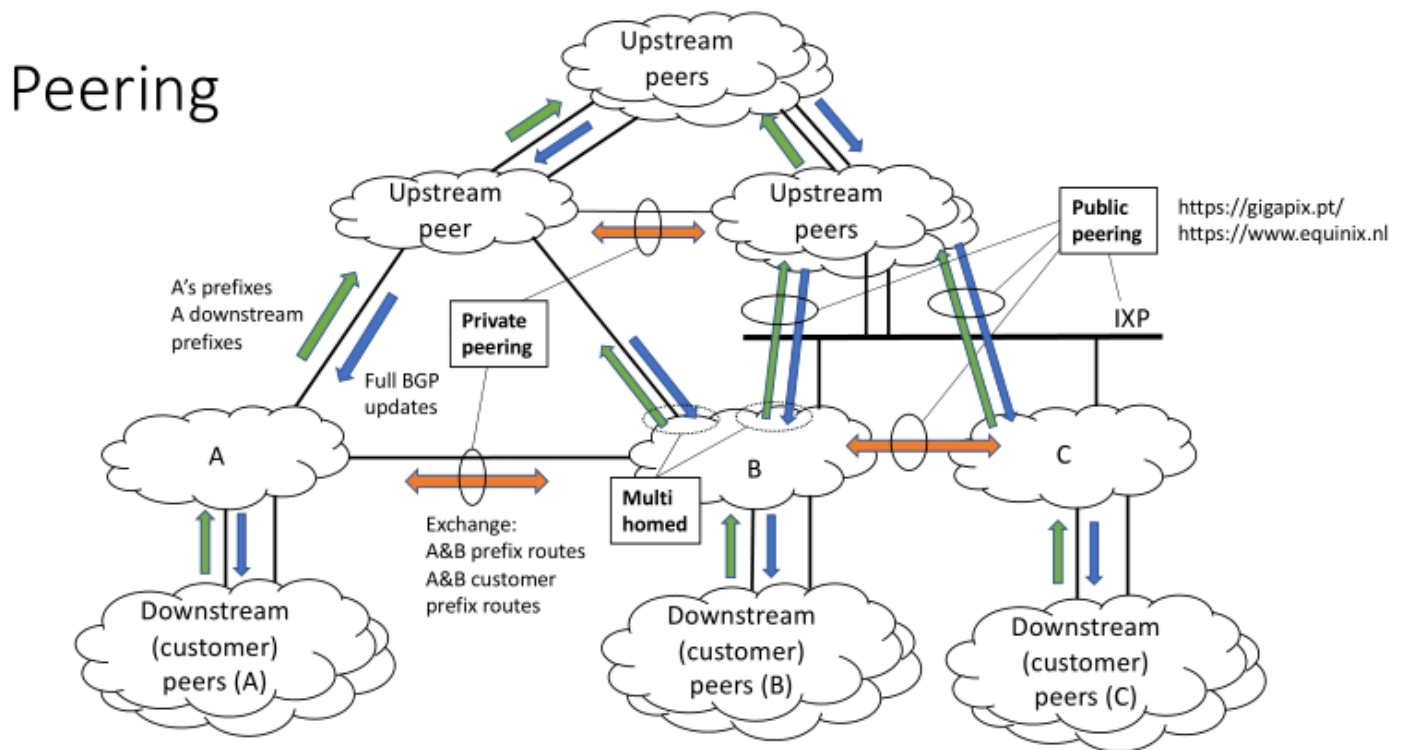
- Different administrative domains:
 - Each node in this network has a different administrator;
 - The Internet;
 - Each node is an Autonomous System - and a network by itself.
- Protocols:
 - **BGP - path vector.**

BGP

- **TCP port 179;**
- Messages:
 - Keep-alive;
 - Update.

- Exchange path-vector routes to other BGP routers:
 - For each network;
 - Select one route, announce it to neighbors, add it to routing table.
- Route map: set of rules to check what to do with route:
 - Drop route, modify route, and add it to the routing table.
- Uses route selection criteria, e.g.:
 - Weight - local to each router;
 - AS-local preference for route - local to each AS;
- Authentication: MD5 hash of password;
- IP header TTL set to 1;
- Route aggregation: 32.32.32.0/24, 32.32.33.0/24 announced as 32.32.32.0/23

Peering



<http://www.bgphelp.com/2017/01/05/bgp-types-of-peering/>

17

Figure 8: peering

Software Defined Networking (SDN)

Switch

- No control, just forwarding.

Default Behaviors	Description
drop all	:(
send to default processing	Switch SDN da cisco pode trazer a stack de software da cisco
encapsular e mandar para o controlador	útil para definir regras on-the-fly (controlador analisa pacote e manda decisão)
forward para o output	:)

Features of SDN

- Hardware abstraction:

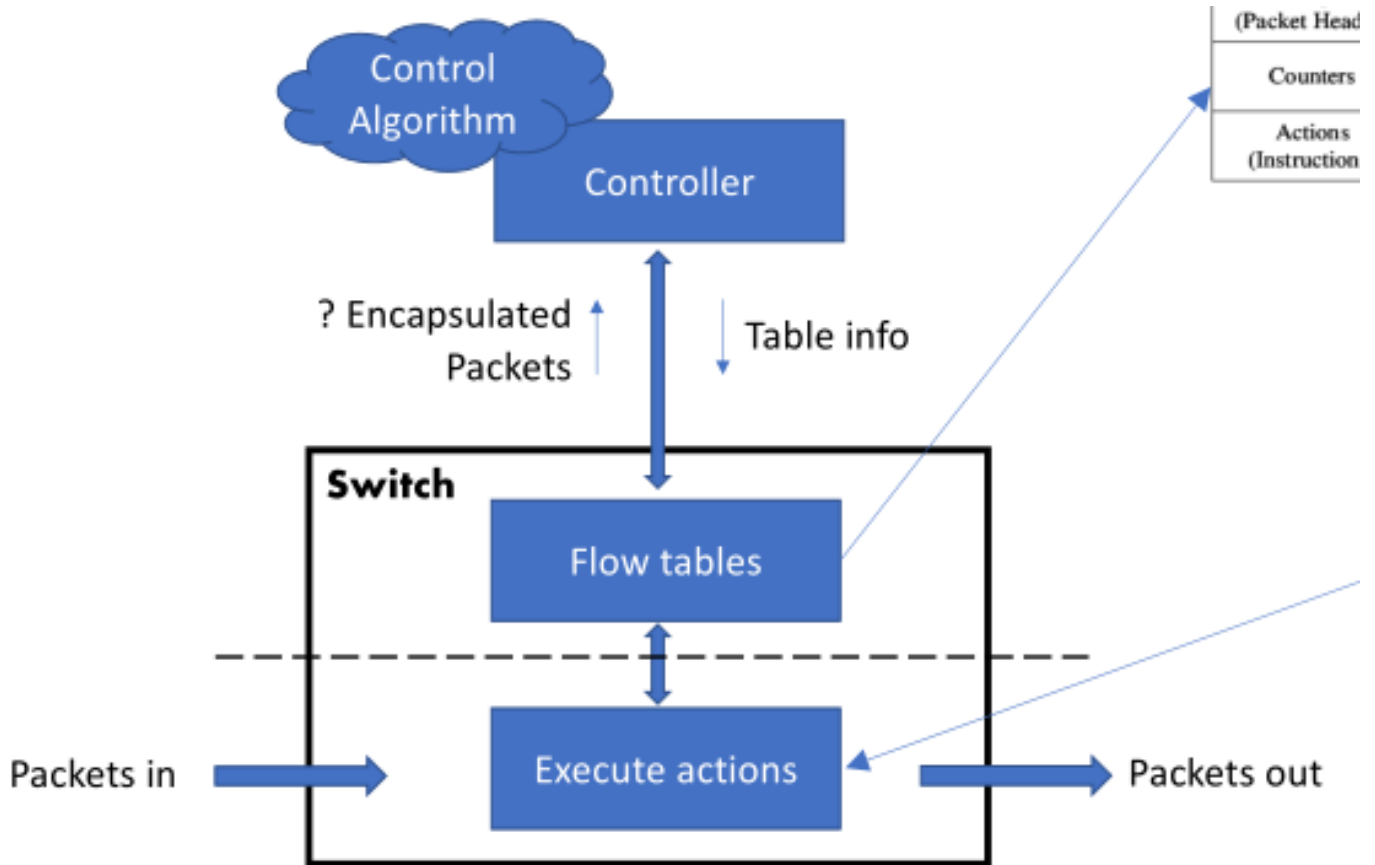


Figure 9: switch_sdn

- Control works on any hardware that implements API.
- Programmable:
 - Hardware/software bundles só permitem config de existing control algorithms;
 - SDN allows us to develop our own algorithms.
- Centralized control of policies:
 - Security policies, e.g., “node of type X can only talk with nodes of type Y”;
 - Routing policies, e.g., “route guest traffic through the firewall”;
 - QoS policies, e.g., “prioritize voice traffic”.

Openflow

- Encapsulado por SSH/TLS;
- Network de gestão em de estar sempre ativa;
- Abstrai as camadas de rede: controlador maneja rules a todos os níveis.

Device API and Automation tools

Building blocks

- module – top of the hierarchy of nodes
- containers – related nodes
- lists – identifies nodes
- leaf – individual attributes of a node
- type – every leaf has an associated type

CI/CD

- Version control - found a problem with code => get back to older version that worked;
- Agile - divide in small parts, implement;

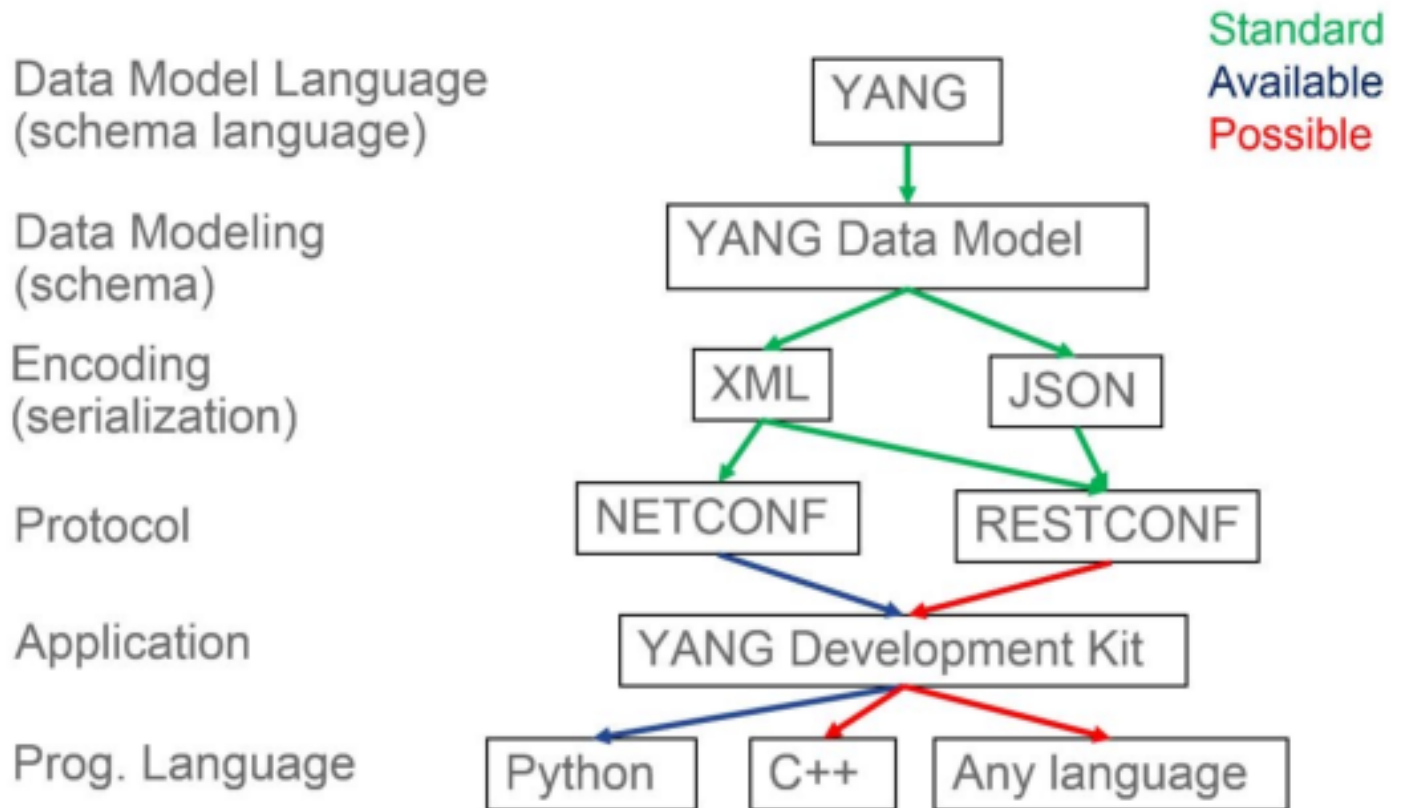


Figure 10: langs

- Pipeline - workflow of code delivery process.

Steps

Types of code

- **Application code** - the software you want to run (including docker/VM image);
- **Infrastructure code** - code that you write to set up the infrastructure;
- **Live infrastructure config** - frontend for infrastructure deployments. Parametrized for a particular deployment.

Code vs. build

- Some languages have a straight forward distinction (C or Java). Source code vs. executable binaries;
- Need a build phase for infrastructure?
 - Depends on the language of the infrastructure code and tools;
 - We actually deploy configuration files.

Test

- **Unit test** - Sandbox or dry run tests of individual components/configurations;
- **System test** - Deploy the different network components and test their interactions;
- **Staging environment** - Replica of production environment.

Release vs. deployment

- **Release** - software product ready to be deployed;
- **Deployment** - release configured for the target environment;
- For infrastructure, the difference is we're dealing with configuration instead of software.



Figure 11: ci_cd_steps

Network Function Virtualization (NFV)

- Não é SDN;
- Pode ser feita usando SDN ou VLANs;

Remote uCPE Architecture




Figure 12: nf_v_sdn

Customer-premises white box (uCPE)

- Run VNFs as VMs or container;
- E.g., DHCP, Routing, NAT, etc.




Figure 13: ucpe